

Mistral 7B

Name

박제현

NLP

2024/05/28


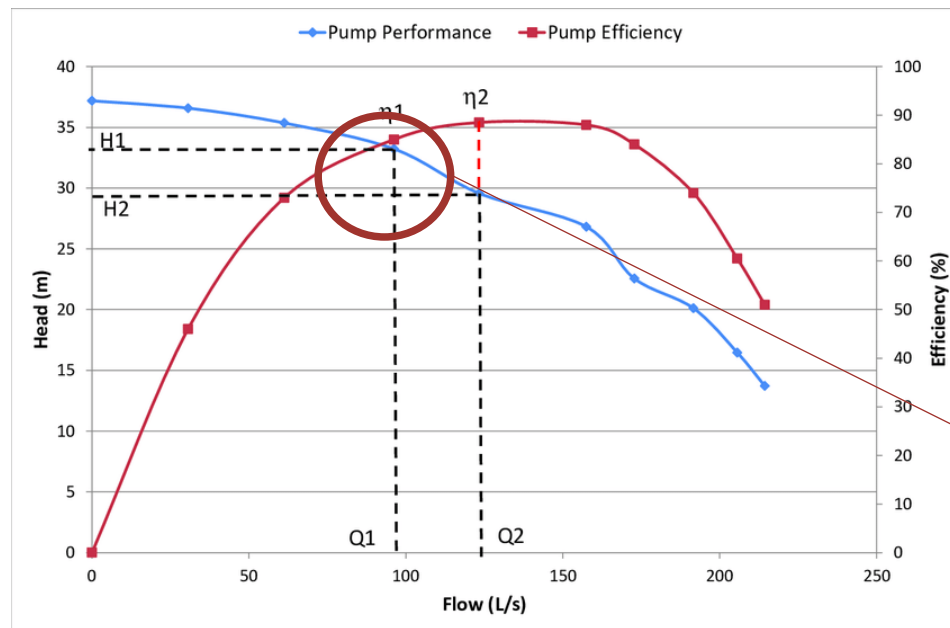


Contents

- Introduction
 - Balancing Performance and Efficiency in NLP Models
 - Recent Approaches – LLMs Optimization
 - Technical Innovations of Mistral 7B
- Architectural details
 - Sliding Window Attention
 - Rolling Buffer Cache
 - Pre-fill and Chunking
- Main results & Analysis

Balancing Performance and Efficiency in NLP Models: Innovations of Mistral 7B

Performance




Model Size & Performance
Performance & 1 / Efficiency

- Lower Computational costs!
- Less Inference latency! -> LoRA

Recent Approaches – LLMs Optimization

Floating point

float 32



1. 32-bit precision requires twice as much GPU memory as **16-bit precision**, allowing more efficient use of GPU memory.
2. Since operations on **lower precision tensors require less memory**, GPUs can process them more quickly.

Recent Approaches – LLMs Optimization

Floating point




Don't transfer all parameters and operations to 16-bit floats. Instead, we **switch between** 32-bit and 16-bit operations during training


Recent Approaches – LLMs Optimization

Floating point

float 16 ("half" precision)



bfloat 16 ("brain" floating point, more "dynamic range" like float 32)




Google developed this format for machine learning and deep learning applications, particularly in their **Tensor Processing Units (TPUs)**. While bfloat16 was originally developed for TPUs, this format is now supported by **several NVIDIA GPUs**.

Recent Approaches – LLMs Optimization

Quantization

Master the Art of Quantization: A Practical Guide with TensorFlow and PyTorch




Quantization converts the model weights from **float32** to **low-bit integer** representations, for example, **8-bit integers** (and, recently, even **4-bit integers**).

Recent Approaches – LLMs Optimization

Quantization

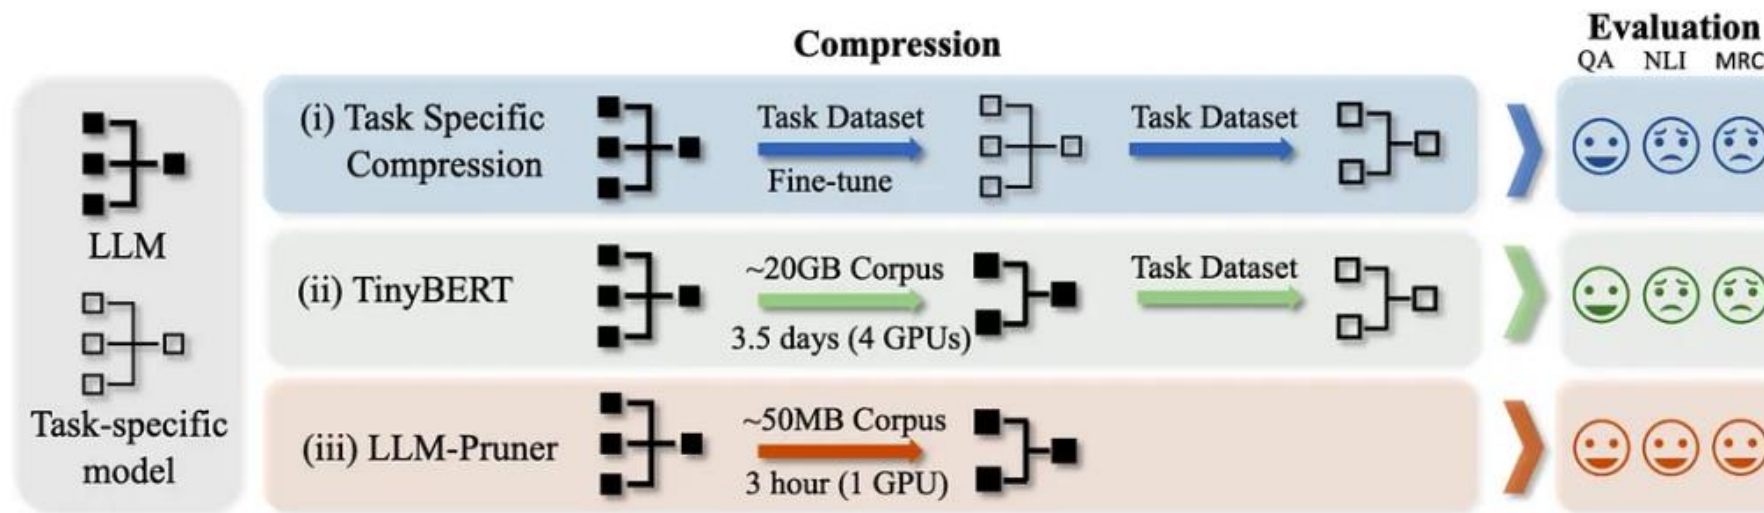
- 1. Post-Training Quantization (PTQ):** A model is **first trained to converge**, then we **convert its weights** to a lower precision without more training. It is usually quite cheap to implement in comparison to training.
- 2. Quantization-Aware Training (QAT):** Quantization is applied **during pre-training or further fine-tuning**. QAT can perform better but requires extra computation resources and access to representative training data

- But Both can lower Performance



Recent Approaches – LLMs Optimization


Pruning



Adopts **structural pruning** that selectively removes **non-critical coupled structures** based on **gradient information**, maximally preserving most of the LLM's functionality.

Recent Approaches – LLMs Optimization


Pruning



Compared to **magnitude pruning** which removes weights **solely based on their magnitudes**, Wanda removes weights on a **per-output basis** by the **product of weight magnitudes and input activation norms**.

Balancing Performance and Efficiency in NLP Models: Innovations of Mistral 7B

Pruning



Compared to **magnitude pruning** which removes weights **solely based on their magnitudes**, Wanda removes weights on a **per-output basis** by the **product of weight magnitudes and input activation norms**.

Technical Innovations of Mistral 7B

Fine-tuning

Grouped-Query Attention (GQA):

- Acceleration of inference speed.
- Reduction of memory requirements and increase in batch size.


Sliding Window Attention (SWA):

- Improved efficiency in handling long sequences.
- Reduction of computational costs.


Mistral AI

Sliding Window Attention


Sliding Window Attention



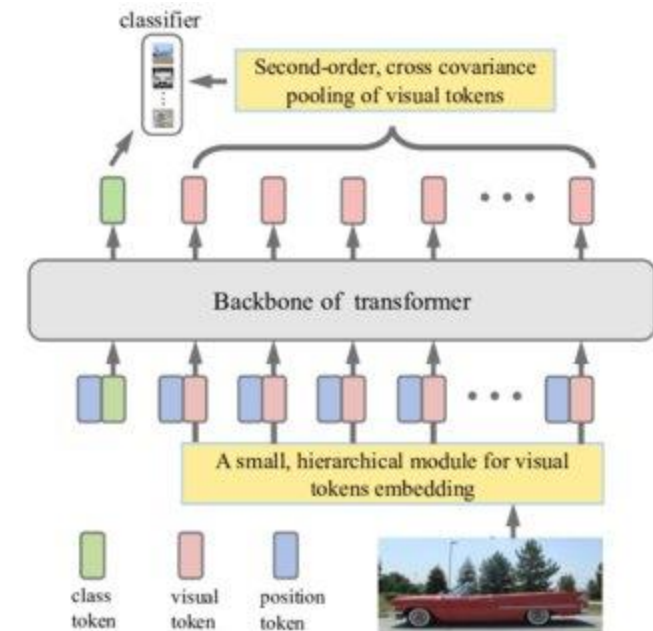
Vanilla Attention



Sliding Window Attention



Effective Context Length



Sliding Window Attention

Sliding Window Attention

Vanilla Attention:

Computational complexity: $O(n^2)$, where n is the sequence length.

Memory usage: Increases **linearly** with the number of tokens.


Doubling the sequence length **quadruples the computation**.

Sliding Window Attention (SWA):

Computational complexity: $O(n * W)$, where W is the window size.

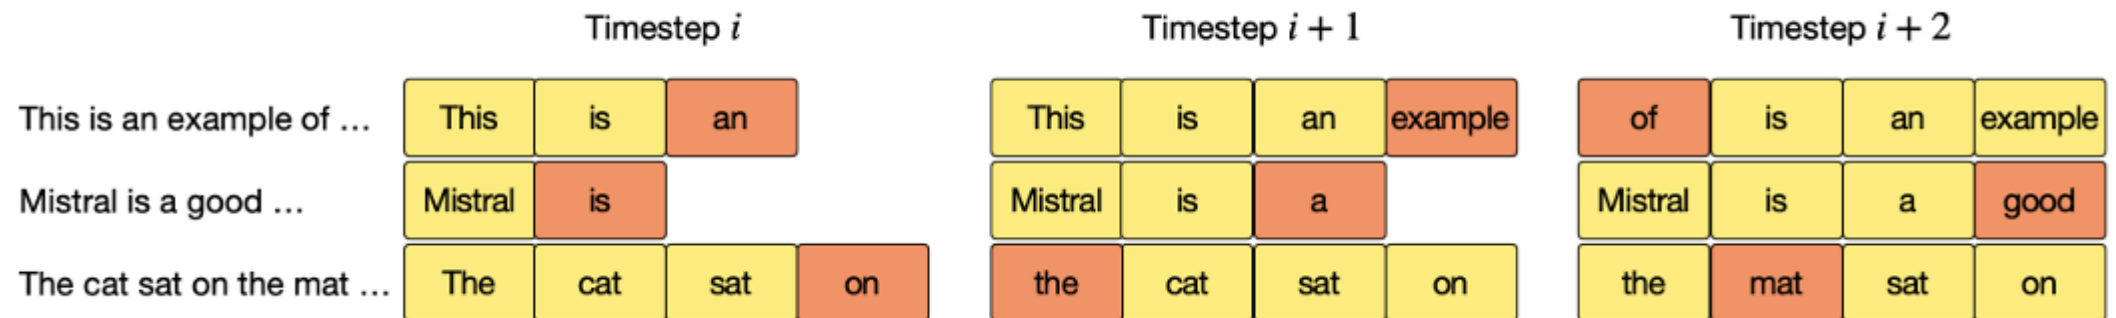
Memory usage: Increases **linearly** with W .

The computational load is **constrained by the window size**, even as the sequence length increases.



Rolling Buffer Cache

Rolling Buffer Cache



The attention span is **limited to a fixed size W** .


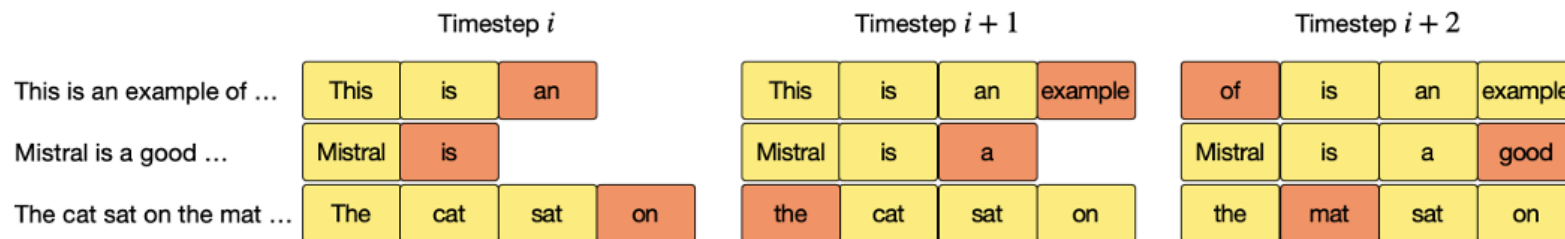
This means each token will only attend to **W preceding tokens**.

The cache is organized as a rolling buffer with a fixed size W .

Keys and values for each timestep i are stored in the cache position $i \bmod W$. ($i = 0, 1, 2, 3, 4 \dots$)

Rolling Buffer Cache

Rolling Buffer Cache



Pre-fill and Chunking



Pre-fill and Chunking

	The	cat	sat	on	the	mat	and	saw	the	dog	go	to
the	0	0	0	0	0	1	1	1	1	0	0	0
dog	0	0	0	0	0	0	1	1	1	1	0	0
go	0	0	0	0	0	0	0	1	1	1	1	0
to	0	0	0	0	0	0	0	0	1	1	1	1
	Past				Cache				Current			

the	mat	and	saw	the	dog	go	to
the	1	1	1	1	dog	go	to
the	mat	and	saw	the	dog	go	to
the	mat	1	1	1	1	go	to
the	mat	and	saw	the	dog	go	to
the	mat	and	1	1	1	1	to
the	mat	and	saw	the	dog	go	to
the	mat	and	saw	1	1	1	1

Main results & Analysis

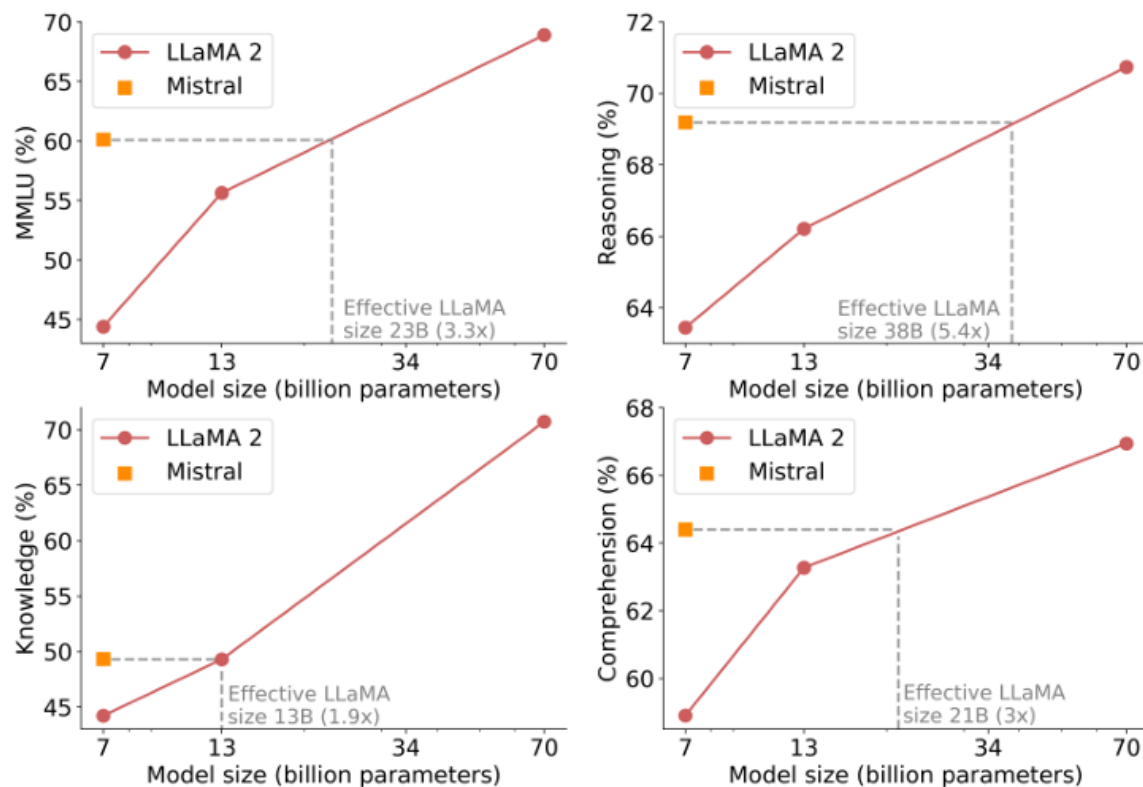
Main results & Analysis



Model	Modality	MMLU	HellaSwag	WinoG	PIQA	Arc-e	Arc-c	NQ	TriviaQA	HumanEval	MBPP	MATH	GSM8K
LLaMA 2 7B	Pretrained	44.4%	77.1%	69.5%	77.9%	68.7%	43.2%	24.7%	63.8%	11.6%	26.1%	3.9%	16.0%
LLaMA 2 13B	Pretrained	55.6%	80.7%	72.9%	80.8%	75.2%	48.8%	29.0%	69.6%	18.9%	35.4%	6.0%	34.3%
Code-Llama 7B	Finetuned	36.9%	62.9%	62.3%	72.8%	59.4%	34.5%	11.0%	34.9%	31.1%	52.5%	5.2%	20.8%
Mistral 7B	Pretrained	60.1%	81.3%	75.3%	83.0%	80.0%	55.5%	28.8%	69.9%	30.5%	47.5%	13.1%	52.2%

Main results & Analysis

Main results & Analysis





TRAIN AND TEST